

“Expert Cloud Consulting”

SOP | Canary Deployment using Docker and NGINX

22 july 2025

—

Contributed by: Akshata

Approved by: Akshay (In Review)

Expert Cloud Consulting

Office #811, Gera Imperium Rise,

Hinjewadi Phase-II Rd, Pune, India – 411057

Canary Deployment using Docker and NGINX






Advanced IaC and Deployment Strategies

Topics :

- Advanced Terraform: modules, workspaces.
- Deployment strategies: Blue-Green, Canary, Rolling updates.

Assignments:

1. **Build a Terraform project with:**
 - Modules for reusable VPC and EC2 resources.
 - Separate workspaces for development and production environments.
2. **Simulate a Canary deployment:**
 - Create two Dockerized versions of a web app.
 - Gradually route traffic to the new version using an Nginx load balancer.



Objective

To simulate a real-world Canary deployment strategy by:

Creating two Dockerized versions of a web application:

v1: Stable and currently live version

v2: New version with potential changes or improvements

Using NGINX as a reverse proxy and load balancer to:

Gradually shift traffic from the stable version (v1) to the new version (v2)

Minimize risk during deployment by allowing controlled exposure of new features

Enable early detection of issues with the new version before full rollout

This setup mimics production-grade deployments where continuous delivery and high availability are crucial, allowing for safe testing of new releases in live environments with real user traffic.

Document References

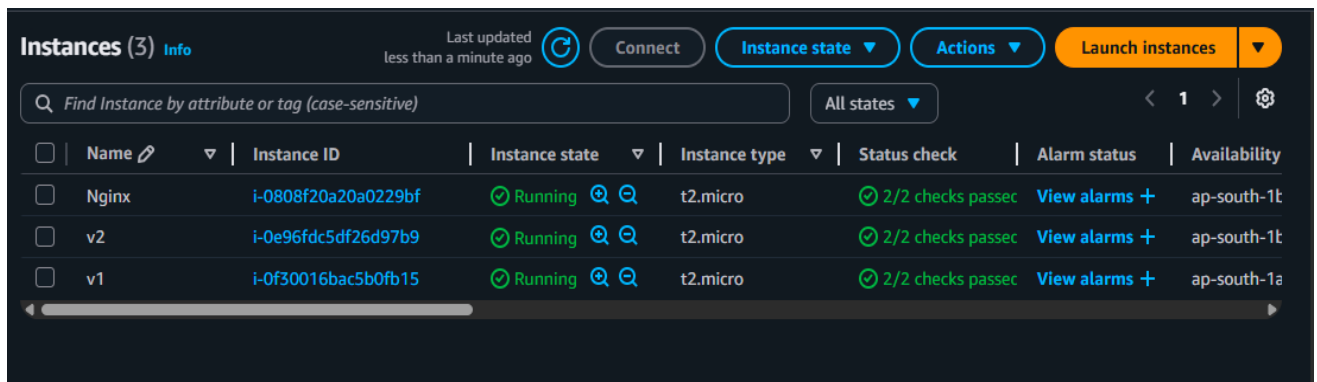
The following resources were referred to during the creation and execution of this Terraform-based infrastructure setup

Date	Document	Filename / Url
17 July	AWS Canary Deployment	https://youtu.be/TzShhyhSpHc?si=U96lHlWVfraMuMWa
18 July	Blue-Green/Canary deployment with NGINX	https://techannotation.wordpress.com/2019/12/13/blue-green-canary-deployment-with-nginx/

To simulate a Canary deployment, I created the following AWS infrastructure using EC2 instances



Total EC2 Instances: 3



	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability
<input type="checkbox"/>	Nginx	i-0808f20a20a0229bf	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1t
<input type="checkbox"/>	v2	i-0e96fdc5df26d97b9	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1t
<input type="checkbox"/>	v1	i-0f30016bac5b0fb15	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1a

1 Public EC2 Instance (Nginx Load Balancer)

Acts as the public-facing load balancer.

Installed with Nginx to route traffic to backend instances.

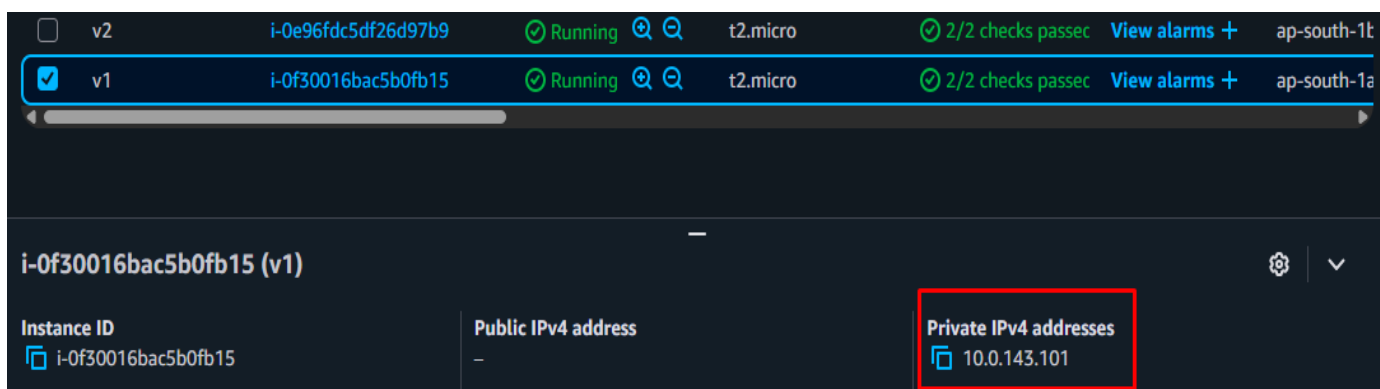
Assigned a public IP.

2 Private EC2 Instances (Application Servers):

Web App v1 Instance

Runs the version 1 of the Dockerized web application.

Accessible only within the VPC via private IP (e.g., 10.0.143.101)



	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability
<input type="checkbox"/>	v2	i-0e96fdc5df26d97b9	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1t
<input checked="" type="checkbox"/>	v1	i-0f30016bac5b0fb15	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1a

i-0f30016bac5b0fb15 (v1)		
Instance ID	Public IPv4 address	Private IPv4 addresses
i-0f30016bac5b0fb15	—	10.0.143.101

Runs the version 2 of the Dockerized web application.

Also private with IP (e.g., 10.0.145.203)



Web App v2 Instance

<input type="checkbox"/>	nginx	i-0808f20a20a0229b1	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1a
<input checked="" type="checkbox"/>	v2	i-0e96fdc5df26d97b9	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1a
<input type="checkbox"/>	v1	i-0f30016bac5b0fb15	Running	t2.micro	2/2 checks passed	View alarms +	ap-south-1a

i-0e96fdc5df26d97b9 (v2)			⚙️	▼
Instance ID	Public IPv4 address	Private IPv4 addresses		
i-0e96fdc5df26d97b9	-	10.0.145.203		

VPC Configuration:

A custom VPC was created with:

2 private subnets (for v1 and v2 app servers)

1 public subnet (for Nginx load balancer)

Route tables and NAT gateway properly configured to allow internet access where needed.

Connecting to Private EC2s (v1 & v2) from NGINX

To securely configure and deploy web applications on private EC2 instances (v1 and v2), SSH access is achieved through the public EC2 instance using its .pem key

```
ubuntu@ip-10-0-23-69:~$ cd .ssh
ubuntu@ip-10-0-23-69:~/ssh$ ls
authorized_keys  aws-key.pem  known_hosts  known_hosts.old
ubuntu@ip-10-0-23-69:~/ssh$
```

```
chmod 400aws-key.pem
```

```
ssh -i "your-key.pem" ubuntu@<v1-private-ip>
```



Once connected to the v1 EC2 instance

```
* Management:      https://landscape.canonical.com
* Support:         https://ubuntu.com/pro

System information as of Tue Jul 22 09:15:44 UTC 2025

System load:  0.08           Processes:            112
Usage of /:   20.5% of 13.49GB Users logged in:          0
Memory usage: 30%           IPv4 address for enx0: 10.0.143.101
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

Failed to connect to https://changelogs.ubuntu.com/meta-release-lts. Check your Internet connection or proxy settings

*** System restart required ***
Last login: Tue Jul 22 06:51:14 2025 from 10.0.23.69
ubuntu@ip-10-0-143-101:~$
```

i-0808f20a20a0229bf (Nginx)

Update the system and install Docker:

```
sudo apt update
sudo apt install docker.io -y
sudo systemctl start docker
sudo systemctl enable docker
```

Create a Dockerfile:

```
FROM nginx:latest
COPY ./index.html /usr/share/nginx/html/index.html
```

"Dockerfile" 2L, 71B

i-0808f20a20a0229bf (Nginx)



Create a simple index.html for v1:

```
<!DOCTYPE html>
<html lang="en">
<head>
<meta charset="UTF-8">
<meta name="viewport" content="width=device-width, initial-scale=1.0">
<title>Web App v1</title>
<style>
body {
font-family: Arial, sans-serif;
text-align: center;
margin-top: 20%;
}
h1 {
color: blue;
}
</style>
</head>
<body>
<h1>Welcome to Web App v1</h1>
</body>
</html>

"index.html" 21L, 357B
```

i-0808f20a20a0229bf (Nginx)

Build and run the Docker container:

docker build -t webapp-v1 .

docker run -d -p 8081:80 --name webapp-v1-container webapp-v1

sudo docker ps

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
webapp	v1	80719b97c490	2 minutes ago	192MB

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
2ab162ded9fa	webapp:v1	"/docker-entrypoint..."	3 seconds ago	Up 3 seconds	0.0.0.0:8081->80/tcp, [::]:8081->80/tcp	mystifying_pike

i-0808f20a20a0229bf (Nginx)

PublicIPs: 43.204.145.51 PrivateIPs: 10.0.23.69

```
*** System restart required ***
Last login: Tue Jul 22 06:51:14 2025 from 10.0.23.69
ubuntu@ip-10-0-143-101:~$ ls
Dockerfile  index.html
ubuntu@ip-10-0-143-101:~$
```

i-0808f20a20a0229bf (Nginx)



connected to the v2 EC2 instance

```
chmod 400aws-key.pem
```

```
ssh -i "your-key.pem" ubuntu@<v2-private-ip>
```

```
Last login: Tue Jul 22 07:35:24 2025 from 10.0.23.69
ubuntu@ip-10-0-145-203:~$ ls
Dockerfile index.html
ubuntu@ip-10-0-145-203:~$
```

i-0808f20a20a0229bf (Nginx)
PublicIPs: 43.204.145.51 PrivateIPs: 10.0.23.69

Update the system and install Docker:

```
sudo apt update
sudo apt install docker.io -y
sudo systemctl start docker
sudo systemctl enable docker
```

Create a Dockerfile:

```
FROM nginx:latest
COPY ./index.html /usr/share/nginx/html/index.htm
```

"Dockerfile" 2L, 71B

i-0808f20a20a0229bf (Nginx)
PublicIPs: 43.204.145.51 PrivateIPs: 10.0.23.69



Create a simple index.html for v2:

```
<!DOCTYPE html>
<html lang="en">
<head>
<meta charset="UTF-8">
<meta name="viewport" content="width=device-width, initial-scale=1.0">
<title>Web App v1</title>
<style>
body {
font-family: Arial, sans-serif;
text-align: center;
margin-top: 20%;
}
h1 {
color: green;
}
</style>
</head>
<body>
<h1>Welcome to Web App v2</h1>
</body>
</html>
```

"index.html" 21L, 358B

i-0808f20a20a0229bf (Nginx)

PublicIPs: 43.204.145.51 PrivateIPs: 10.0.23.69

Build & Run Docker container

`docker build -t webapp-v2 .`

`docker run -d -p 8081:80 webapp-v2`

`sudo docker ps`

docker: Error response from daemon: pull access denied for webapp, repository does not exist or may require 'docker login': denied: requested access to the resource is denied

Run 'docker run --help' for more information

ubuntu@ip-10-0-145-203:~\$ sudo docker run -d -p 8082:80 webapp:v1

Unable to find image 'webapp:v1' locally

docker: Error response from daemon: pull access denied for webapp, repository does not exist or may require 'docker login': denied: requested access to the resource is denied

Run 'docker run --help' for more information

ubuntu@ip-10-0-145-203:~\$ sudo docker run -d -p 8082:80 webapp:v2

ad8affdbe6086c711a59796673452d2cccf74c0fc585e88580c81fbd307b0bb6

ubuntu@ip-10-0-145-203:~\$ sudo docker ps

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
ad8affdbe608	webapp:v2	"/docker-entrypoint..."	6 seconds ago	Up 5 seconds	0.0.0.0:8082->80/tcp, [::]:8082->80/tcp	beautiful_rosalind

ubuntu@ip-10-0-145-203:~\$

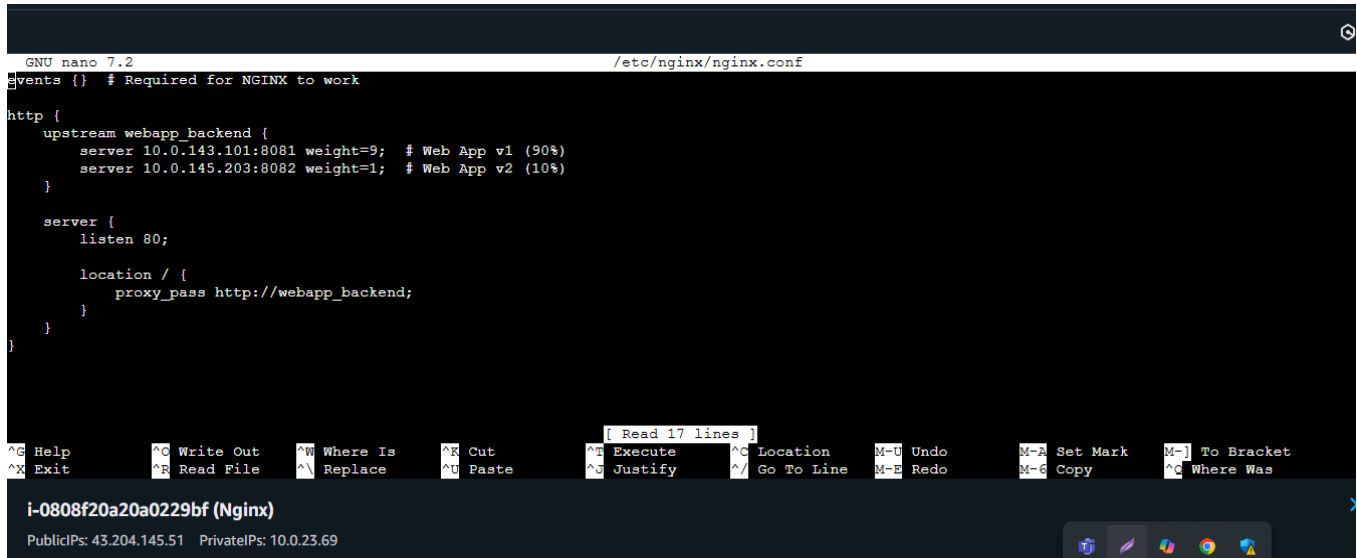
i-0808f20a20a0229bf (Nginx)

PublicIPs: 43.204.145.51 PrivateIPs: 10.0.23.69



Nginx Configuration on Public EC2

/etc/nginx/nginx.conf:



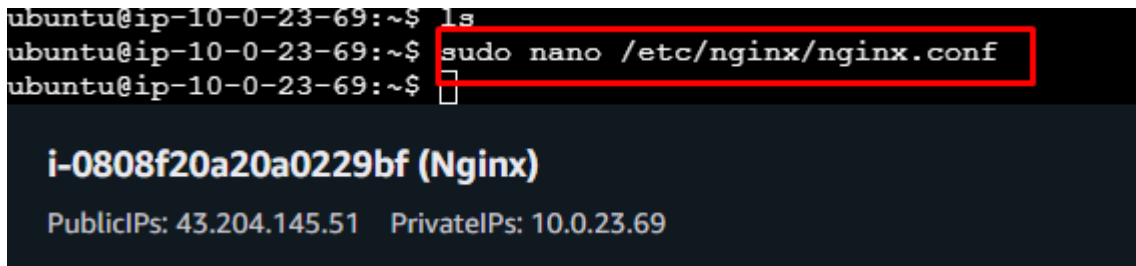
```
GNU nano 7.2 /etc/nginx/nginx.conf
events {} # Required for NGINX to work

http {
    upstream webapp_backend {
        server 10.0.143.101:8081 weight=9; # Web App v1 (90%)
        server 10.0.145.203:8082 weight=1; # Web App v2 (10%)
    }

    server {
        listen 80;

        location / {
            proxy_pass http://webapp_backend;
        }
    }
}
```

ls



```
ubuntu@ip-10-0-23-69:~$ ls
ubuntu@ip-10-0-23-69:~$ sudo nano /etc/nginx/nginx.conf
ubuntu@ip-10-0-23-69:~$
```

i-0808f20a20a0229bf (Nginx)
PublicIPs: 43.204.145.51 PrivateIPs: 10.0.23.69


Restart nginx:

```
sudo nginx -t
sudo systemctl restart nginx
```

Open Browser and Paste Public IP

Now, go to your browser and open:

<http://43.204.145.51>



Welcome to Web App v1



Welcome to Web App v2

